



In one line: Every structure is a **memory layout** plus an **invariant** that makes some operations cheap and others expensive. Pick by the operation that must be fast, then check the constants: a trip to main memory costs $\sim 50\times$ an L1 hit, so **contiguous memory + cache lines** beat pointer chasing at equal big-O.


Array: append past capacity




Linked list: pointer chasing



Hash table: linear probing



Min-heap = complete tree in array



How each one works

- Array** — one block, element i at base + $i \cdot \text{stride}$ (MemoryLayout<T>.stride) $\Rightarrow O(1)$ index. Scans ride the prefetcher; a cache line (64 or 128 B) brings 8–16 Ints per miss. Middle insert / delete shifts $O(n)$. [SomeStruct] holds values inline; [SomeClass], a Python list, a Java ArrayList<Integer> hold *references* — contiguous pointers to scattered objects.
- Linked list** — one allocation per node; a textbook doubly linked list pays two links per datum (200 % overhead) and one malloc each — why CPython's deque links 64-slot blocks instead. Wins only for $O(1)$ splice at a *held* node and stable addresses; finding it is $O(n)$.
- Hash table** — hash(key) mod $m \rightarrow$ bucket, load $\alpha = n/m$; chaining (Java HashMap bins) or open addressing (Swift, Python).
- BST** — left < node < right, $O(h)$; sorted inserts give $h = n$. **AVL**: subtree heights differ ≤ 1 , $h \leq 1.44 \log_2 n$, faster lookups. **Red-black**: no red-red, equal black height, $h < 2 \log_2(n+1)$, ≤ 2 rotations per insert, ≤ 3 per delete (Java TreeMap). **B-tree**: one node = one disk page (SQLite: 4096 B); fan-out $f \Rightarrow$ height $\log_f n$ (100-way: 10^6 keys in 3 levels).
- Skip list** — a sorted linked list plus random “express lanes”: expected $O(\log n)$ with no rebalancing; Java's ConcurrentSkipListMap lets threads insert and remove concurrently; Redis sorted sets = skip list + hash table.
- Binary heap** — complete tree in an array, parent $\leq$ children, peek $O(1)$; push = append + sift-up, pop = last to root + sift-down, $O(\log n)$; **heapify $O(n)$** . Swift's Heap is **min-max**: even levels $<$ descendants, odd levels $\geq \Rightarrow$ both min and max $O(1)$.
- Trye** — node per prefix: $O(L)$ per op whatever n is; answers prefix queries a hash set can't. A **hash trie** (CHAMP: Swift TreeDictionary) walks 5 hash bits per level into 32-way nodes; copies share every unchanged subtree.
- Graph** — adjacency list $O(V+E)$ space, neighbours $O(\deg)$; matrix $V \times V$ ($V^2/8$ bytes as bits), “edge $u \sim v$?” $O(1)$, neighbours $O(V)$. BFS / DFS $O(V+E)$ on a list; Dijkstra with a binary heap $O((V+E) \log V)$.
- Union-find** — parent array + path compression + union by rank $\Rightarrow O(\alpha(n))$, $\alpha \leq 4$ for any real n : components, Kruskal.

Big-O — average (worst)

structure	access	search	insert / delete	note
array	1	n	n ; end: 1^{am}	
sorted array	1	$\log n$	n	the shift dominates
linked list	n	n	1 at node	find the node: n
stack/queue/deque	ends 1	n	ends 1	array queue front: n
hash table	—	1 (n)	1^{am} (n)	hashing key = $O(\text{len})$
BST, plain	—	$\log n$ (n)	$\log n$ (n)	sorted input $\rightarrow n$
AVL / red-black	—	$\log n$	$\log n$	in-order = sorted
skip list	—	$\log n^{\text{exp}}$	$\log n^{\text{exp}}$	randomised levels
binary heap	min 1	n	$\log n$	build $O(n)$
trie	—	L	L	L = key length
adj. list / matrix	—	$\deg / 1$	$1 / 1$	space $V+E / V^2$

$n = 10^6$; $\log_2 n \approx 20$; $n \log n \approx 2 \times 10^7$; $n^2 = 10^{12}$ — at 10^9 steps/s: 20 ns · 20 ms · 17 min.

Same roles, three runtimes

role	Swift	Java / Kotlin (JVM)	Python
growable array	Array	ArrayList	list
hash map	Dictionary	HashMap = hashMapOf()	dict
insertion order	OrderedDictionary ^c	LinkedHashMap = Kotlin mutableMapOf()	dict (guaranteed since 3.7)
sorted map	SortedDictionary ^c (B-tree, unstable trait)	TreeMap (red-black) = sortedMapOf()	list + bisect
queue / stack	Deque ^c · Array	ArrayDeque (not legacy Stack)	collections.deque
priority queue	Heap ^c (min-max)	PriorityQueue (min)	heapq (min; max 3.14)
binary search	partitioningIndex ^a	Collections.binarySearch	bisect


^c swift-collections · ^a swift-algorithms — the Swift stdlib itself has no deque, heap or sorted tree.

Growth factors — code, not contract

container	new capacity when full (source code)
Swift Array, Rust Vec	$\times 2$
Java ArrayList	$\times 1.5$ (old + (old >> 1))
Java ArrayDeque	$\times 2$ below 64, then $\times 1.5$
Go slice	$\times 2$ below 256, then easing to $\times 1.25$
Python list	$n + n/8 + 6$, rounded to 4: 0, 4, 8, 16, 24, 32, 40, 52, 64, 76 ...

Apple's and Java's docs promise only amortised-constant append — the factor is an implementation detail. $2\times$ can never reuse the blocks it freed (their sum is always smaller than the next request); $1.5\times$ can after 4 reallocations (folly's fbvector rationale).

One graph, two layouts



Example — freeing a long class list


```
final class Node<T> {
  var value: T; var next: Node?
  init(_ v: T, next: Node? = nil) { value = v; self.next = next }
  deinit { // default: each deinit frees the next → recursion
    var node = self // walk instead, while the link is unique
    while isKnownUniquelyReferenced(&node.next) {
      (node, node.next) = (node.next!, nil)
    }
  }
}
```

The loop stops at nil (false) or at the first node another owner still holds — that node and its tail stay alive.

Inside the Swift types

- Set / Dictionary**: 2^k buckets, $\leq 3/4$ full, an occupancy bitmap. Seed = the table's address; new on every allocation and resize (kept for COW copies) — two equal sets need not iterate alike, even in one run.
- Three ways to keep insertion order**: OrderedSet and Python's dict = dense array + a hash table of indices (middle insert / remove renumbers: $O(n)$ in OrderedSet); Java's LinkedHashMap threads a doubly linked list through its entries (access order $\rightarrow$ an LRU cache).
- Deque**: circular buffer (may wrap). A class-element Array may be an NSArray, ContiguousArray never. InlineArray (26.0): no own allocation, copied eagerly.

Layout beats big-O: binary search



Traps

- A 500k-node class list overflows the stack** in its default deinit (indirect enums: no source fix). A strong prev is a retain cycle — weak.
- “Linked lists insert faster”** — only at a held node; the $O(n)$ walk and cache misses lose to Array's memmove.
- Amortised $\neq$ worst case**: one append can copy $O(n)$ inside a 16 ms frame or an audio callback — reserveCapacity, init(minimumCapacity:), a fixed ring.
- A heap is not sorted**: contains / remove(x) are $O(n)$, and Dijkstra's decrease-key needs a vertex $\rightarrow$ position map beside it.

Remember

Layout decides speed, invariant decides big-O. Contiguous by default; pointers only when you must splice.