


In one line: Every schema change must be tolerated by **all code that is still running** — the old pods of a rolling deploy, the blue stack you may roll back to, and **app versions installed on phones for months**. So a breaking change becomes **expand** → **migrate** → **contract** across several releases, data moves in **small batches**, and every DDL holds its lock for **milliseconds**.



Why every breaking change takes several releases

- A **rolling deploy** runs N and N+1 against one DB; a **blue/green rollback** brings N back; **mobile clients** are version N-k for months. A rename in one step breaks at least one of them.
- Rule: schema *v* must work with code N and N+1; code N+1 must work with schema *v* and *v*+1.
- Same for the API: add the field, keep the old one until its clients are gone; for app clients measure by version header, then raise the min version.

Expand / contract — rename name → full_name

- Expand**: add the column nullable (or constant default), tables, indexes. Old code ignores them.
- Dual write** (N+1): write both, still read old. Or a DB trigger — keeps writes from old code in sync too.
- Backfill** old rows in batches (below). Verify: counts, checksums, sample diffs.
- Switch reads** (N+2); then constraints: NOT NULL, FK, unique.
- Stop writing** the old column; make the ORM ignore it (Rails `ignored_columns`) or its cached column list breaks on drop.
- Contract**: drop column / trigger — after rollback window + old clients.

Backfill — batches, not one UPDATE

One UPDATE of 50M rows = one huge transaction: row locks held for its whole length, WAL burst, replica lag, dead tuples. Instead: keyset batches (1-10k rows by PK), commit each, sleep, throttle on replica lag; **idempotent** (WHERE full_name IS NULL) and **resumable**; run as a job, not inside the deploy.

MySQL

ALGORITHM=INSTANT ADD/DROP COLUMN (default in 8.4, any position; 64 row versions, then a rebuild). Other changes: **gh-ost** — triggerless: copies into a ghost table and applies the **binlog** stream; throttle / pause, cut-over when you choose; **pt-online-schema-change** — same copy, but **triggers** on the live table (extra write load). Both end with an atomic RENAME.

Dual writes vs CDC (moving data to a new store)

App **dual write** = two non-atomic writes: one fails or reorders ⇒ silent drift; needs reconciliation. **CDC** (Debezium-style reading WAL/binlog) or a **transactional outbox**: one source of truth, ordered by log position, replayable. Snapshot + stream → verify → switch reads → switch writes.

Order around deploys · rollback

Expand before the code that needs it (old code must survive it); **contract after** the old code and the rollback window are gone. One migrator with an advisory lock — not every pod at start-up. A “down” migration after data is written loses data: roll **forward**; contract is the only one-way door, so back up (PITR) before it.

PostgreSQL — which lock, which trick

statement	lock · safe form
most ALTER TABLE	ACCESS EXCLUSIVE : blocks reads <i>and</i> writes — must be instant
ADD COLUMN ... DEFAULT	PG 11+: constant default kept in metadata, no rewrite; <i>volatile</i> (clock_timestamp()) rewrites rewrite + index rebuild, unless binary coercible (varchar(n) → text)
ALTER COLUMN TYPE	SHARE: blocks writes → CONCURRENTLY
CREATE INDEX	NOT VALID (instant) then VALIDATE CONSTRAINT (SHARE UPDATE EXCLUSIVE: reads + writes go on). FK locks <i>both</i> tables
ADD CONSTRAINT ... CHECK / FOREIGN KEY	full scan — unless a valid CHECK (c IS NOT NULL) proves it. PG 18: NOT NULL itself can be NOT VALID
SET NOT NULL	build the unique index CONCURRENTLY, then ADD CONSTRAINT ... USING INDEX

Example — PostgreSQL, one step per deploy

```
SET lock_timeout = '2s'; -- fail fast + retry, never queue
ALTER TABLE users ADD COLUMN full_name text; -- metadata only
-- deploy N+1 (writes both); then repeat until 0 rows:
UPDATE users SET full_name = name WHERE id IN (SELECT id
  FROM users WHERE full_name IS NULL ORDER BY id LIMIT 5000);
ALTER TABLE users ADD CONSTRAINT full_name_nn
  CHECK (full_name IS NOT NULL) NOT VALID; -- instant
ALTER TABLE users VALIDATE CONSTRAINT full_name_nn; -- no write lock
ALTER TABLE users ALTER full_name SET NOT NULL; -- CHECK proves it
CREATE INDEX CONCURRENTLY users_full_name_idx ON users (full_name);
-- deploy N+2; weeks later, when N, N+1 and app 5.1 are gone:
ALTER TABLE users DROP COLUMN name; -- contract
```

The lock queue — how a 1 ms ALTER causes an outage

The ALTER waits for ACCESS EXCLUSIVE behind one long SELECT or an *idle in transaction* session; every query that arrives *after* it queues behind the ALTER. Lock_timeout applies to each lock attempt: give up after 1-5 s, retry with backoff. CONCURRENTLY: two scans, waits for old transactions, not inside a transaction block (turn off the tool's wrapping txn), failure leaves an **INVALID** index — drop it and retry, or REINDEX INDEX CONCURRENTLY.

Traps

- ! Rename / drop in the same release as the code change: old pods and old apps crash.
- ! New NOT NULL column with no default: old code's INSERTs fail.
- ! “The ALTER is instant” — not while it waits in the lock queue.

Remember

Add, write both, backfill, read new, stop old, drop — each its own deploy.