



In one line: A CIImage is a **recipe, not pixels**: applying a CIFilter only extends a lazy graph. Nothing is computed until a CIContext **renders** a rectangle of it — then Core Image walks the graph *backwards* from that rectangle (region of interest), **concatenates** adjacent kernels into as few GPU (Metal) passes as it can, and writes the result to a CGImage, pixel buffer, texture or file. Contexts are **expensive** — **make one and reuse it**; contexts and images are immutable and thread-safe, filters are mutable and are not.

A. Build a recipe (free) → render a rectangle (all the work)



B. Extent — the blur trap



Outside its extent an image is clear black. Clamp (repeat the edge pixels, infinite extent) *before* the blur, crop back to input.extent *after*. Generators (CIColorGenerator, checkerboard) are infinite too: rendering .infinite fails — always crop.

How it works

- Inputs:** CIImage(cgImage:), (cvPixelBuffer:) (camera frames, no copy), (image:) from UIImage — which *drops* imageOrientation: apply .oriented(_): Origin is **bottom-left**, like Quartz.
- Two filter APIs:** stringly CIFilter(name: "CIGaussianBlur") + setValue(_:forKey: kCIInputImageKey) (typos = runtime nil), or typed import CoreImage.CIFilterBuiltins (iOS 13) → CIFilter.gaussianBlur() with inputImage / radius properties.
- Lazy + concatenated:** chains of per-pixel (colour) kernels fuse into one program — no intermediate texture per step; a blur or warp needs neighbours, so it may force one.
- CIContext** owns an MTLCommandQueue, the compiled-kernel cache and intermediate buffers. Apple: *one per view that renders, or per background task* — never per image or per frame. CIContext(mtlDevice:) / (mtlCommandQueue:) shares your Metal pipeline; .useSoftwareRenderer: true forces the (slow) CPU path.
- Colour:** inputs are converted into the **working space** — default **extended sRGB, linear gamma**, so maths happens in linear light and values beyond 0–1 survive; output converts to the destination's space. Masks, depth, data: .workingColorSpace: NSNull() turns colour management off.
- Custom kernels** are Metal (MSL): CIColorKernel (one pixel in, one out), CIWarpKernel (returns the *source coordinate*), CIBlendKernel, general CIKernel (samplers + an ROI callback). Build: *Other Metal Compiler Flags* -fcikernel, user setting MTLINKER_FLAGS = -cikernel. String kernels (CIKernel(source:)) are **deprecated since iOS 12**. MPS or a model inside the graph: CIImageProcessorKernel.
- Vision hand-off:** the same CIImage feeds VNImageRequestHandler(ciImage:orientation:); results are **normalised** (0–1), **bottom-left** origin — convert with VNImageRectForNormalizedRect, then e.g. pixellate only the face rect and composite it back.

Example — one shared context, typed filters

```
import CoreImage.CIFilterBuiltins
let context = CIContext() // ONE, app-wide or per view
func stylise(_ cg: CGImage) -> CGImage? {
    let input = CIImage(cgImage: cg)
    let blur = CIFilter.gaussianBlur() // typed API, iOS 13
    blur.inputImage = input.clampedToExtent() // no fading edges
    blur.radius = 8
    let tone = CIFilter.sepiaTone() // filters: NOT thread-safe
    tone.inputImage = blur.outputImage; tone.intensity = 0.6
    guard let out = tone.outputImage?.cropped(to: input.extent)
    else { return nil } // still only a recipe
    return context.createCGImage(out, from: out.extent) // GPU work: HERE
}
```

Custom colour kernel (MSL) + applying it

```
#include <CoreImage/CoreImage.h> // Threshold.ci.metal
extern "C" float4 threshold(coreimage::sample_t s, float t) {
    float l = dot(s.rgb, float3(0.2126, 0.7152, 0.0722)); // luma
    return float4(float3(step(t, l)), s.a); // black or white
} // Swift side:
let k = try CIColorKernel(functionName: "threshold",
                           fromMetalLibraryData: metalLibData)
let bw = k.apply(extent: img.extent, arguments: [img, 0.5])
```

Live camera — stay on the GPU

AVCaptureVideoDataOutput → CVPixelBuffer → CIImage(cvPixelBuffer:) → filters → context.render into the MTKView drawable's texture (or a CIRenderDestination) — never createCGImage per frame: that is a GPU→CPU readback plus an allocation, 30–60 times a second.

Which tool for an image effect?

Core Image	built-ins, graph optimisation, photo + video, RAW
SwiftUI shader	one view's pixels (.colorEffect / .layerEffect, iOS 17)
Metal	full control: your own passes, compute, 3D, ML
Core Graphics	drawing shapes and text, not filtering pixels

Traps

- ! A new CIContext per image or per frame — kernel recompiles, memory spikes, stutter.
- ! UIImage(ciImage:) in a UIImageView: still a recipe, rendered on every draw by a context you do not control. Render to a CGImage with your shared context.
- ! Soft transparent borders after a blur, or createCGImage(out, from: out.extent) is nil (infinite extent) — clamp, then crop.
- ! Rotated photo — CIImage(image:) drops the orientation.
- ! One CIFilter shared between threads — a data race; share images and the context instead.
- ! Face box drawn in the wrong place — Vision is normalised and bottom-left, UIKit top-left.

Remember

"CIImage is the recipe, CIContext is the kitchen — build the kitchen once, cook only what's ordered."