

In one line: GraphQL is a typed query language over **one endpoint**: the server publishes a **schema**, the client sends an **operation** naming exactly the fields it wants (+ typed **variables**), and gets back JSON of the same shape as `{data, errors}`. **Apollo iOS** generates Swift types from your `.graphql` files and keeps a **normalised cache**: every object with a cache ID is stored *once*, so every query that contains it sees the same, updated value.

1 · response = a tree

```
# query Hero($id: ID!)
{
  hero: {
    __typename: "Droid",
    id: "2001", "name": "R2-D2",
    friends: [
      { __typename: "Human",
        id: "1000", "name": "Luke"},
      { __typename: "Human",
        id: "1003", "name": "Leia"}]
  }
}
```

2 · flattened records, keyed by cache ID

```
QUERY_ROOT
hero(id:"2001") → Droid:2001

Droid:2001
name: "R2-D2"
friends: [Human:1000, Human:1003]

Human:1000
name: "Luke"
→ "Luke S."

Human:1003
name: "Leia"
```

key = `__typename:keyFields`, from `@typePolicy(keyFields: "id")`

no cache ID → keyed by path (`QUERY_ROOT.hero.friends.0`):
a second query gets its own copy — nothing is shared

3 · one record, many screens

```
query Profile {
  human(id: "1000") {
    id name
  }
}

mutation Rename {
  rename(id: "1000",
    to: "Luke S.") {
    id name
  }
}
```

the mutation selects id → its result
lands in **Human:1000** → **both**
watchers (Hero, Profile) re-emit, no refetch

GraphQL

- Schema (SDL)**: object types, scalars (`Int` = 32-bit, `Float`, `String`, `Boolean`, `ID`), `!` = non-null, lists, enums, **interfaces / unions**, input types. Introspection lets tools download it.
- Operations** over `POST /graphql` with `{query, operationName, variables}`: **query** (read; sibling fields may resolve in parallel), **mutation** (write; top-level fields run in order), **subscription** (a stream of events, usually over WebSocket).
- Fragments** = named, reusable selection sets (fragment X on Character); **inline fragments** ... on Droid pick fields per concrete type of an interface / union — which is why the client asks for `__typename`.
- Variables** are typed (`$id: ID!`) and sent separately — never string-build a query.
- vs REST**: no **over-fetching** (only the fields asked for), no **under-fetching** (nested data in one round trip, not N calls), a typed contract. Costs: one URL + POST → no free HTTP/CDN caching; query-cost limits; schema evolution by *deprecating* fields, not versioned URLs.
- Server N+1**: one resolver per field → friends for 50 heroes = 50 DB queries. Fix: **DataLoader** — batch the keys requested in one tick into one query, cache per request.
- Errors**: `"data": { ... }, "errors": [{message, locations, path, extensions}]`. A failing field becomes null (bubbling up to the nearest nullable parent) and the rest is **partial data**. Typically **HTTP 200** anyway — check errors, not the status.

Example — operation + Apollo iOS 2.x call

```
# Hero.graphql (codegen -> HeroQuery, HeroDetails)
fragment HeroDetails on Character { id name }
query Hero($id: ID!) {
  hero(id: $id) { ...HeroDetails friends { ...HeroDetails } }
}
# schema extension (.graphqls): cache ID per type
extend type Human @typePolicy(keyFields: "id")

let store = ApolloStore(cache: try SQLiteNormalizedCache(fileURL: url))
let client = ApolloClient(networkTransport: transport, store: store)
let res = try await client.fetch(query: HeroQuery(id: "2001"),
  cachePolicy: .cacheFirst)
if let errors = res.errors { log(errors) } // GraphQL errors: not thrown
show(res.data?.hero?.name) // may be partial
// network / parsing failure -> thrown; catch it
```

Traps

- ! "200 OK, so it worked" — read `errors[]`; data may be partial or null.
- ! Not selecting `id` (or no `@typePolicy`) — path keys, duplicate copies, a mutation result that updates nothing on screen.
- ! Blaming GraphQL for N+1 — it is the resolvers; DataLoader batches.
- ! "GraphQL is cached like REST" — POST to one URL; the client cache does the work (or persisted-query GETs).
- ! Forgetting to `cancel()` a watcher; one giant query instead of per-view fragments.
- ! Expecting Apollo iOS to roll back an optimistic write for you — it has no `optimisticResponse`; you revert.
- ! 2.x: GraphQL Int *inputs* are generated as Swift `Int32` (the spec's 32 bits).

Remember

Ask for the shape, get the shape; flatten by `__typename:id`; 200 can still be an error.

Apollo iOS — the moving parts

- Codegen**: `apollo-ios-cli (init, fetch-schema, generate)` + `apollo-codegen-config.json`. Schema + `.graphql` operations → a schema-types module + one Swift type per operation / fragment, with nested Data models. Fields you didn't select don't exist in Swift.
- 2.x = async/await rewrite**, Swift 6 strict concurrency, most types Sendable; iOS 15+; CocoaPods dropped. `fetch` returns a `GraphQLResponse (data + errors)`; multi-result policies return an `AsyncThrowingStream`. Interceptors split into `GraphQLInterceptor`, `HTTPInterceptor` (e.g. auth header), `CacheInterceptor`, `ResponseParsingInterceptor`. A cache *miss* is nil, not a thrown error.
- Cache policies** (2.x): `.cacheFirst` (default) · `.networkFirst` (network, cache on failure) · `.networkOnly` · `.cacheOnly` · `.cacheAndNetwork` (cached, then fresh: two emissions). 1.x names: `returnCacheDataElseFetch`, `fetchIgnoringCacheData`, ...
- Cache**: `ApolloStore` over `InMemoryNormalizedCache` or `SQLiteNormalizedCache` (persists). Cache IDs via `@typePolicy(keyFields:)` (`@fieldPolicy` maps a field's arguments to an ID) or programmatically in `SchemaConfiguration`.
- Watchers**: `client.watch(query:)` → `GraphQLQueryWatcher`; re-delivers whenever *any* record it read changes. `cancel()` it when the owner goes — or it leaks.
- Optimistic updates**: Apollo iOS has **no** `optimisticResponse` (Apollo Client web and Kotlin do). Write the expected value yourself in `store.withinReadWriteTransaction` (a `@apollo_client_ios_localCacheMutation` operation gives mutable models), then send the mutation; on error write the old value back.
- Pagination**: Relay **connections** — `first/after` args, `edges { cursor node }`, `pageInfo { hasNextPage endCursor }`. The `apollo-ios-pagination` package's `GraphQLQueryPager` (`fetch`, `loadNext`, `loadPrevious`, `refetch`) watches and merges pages; cursor or offset.
- Persisted queries**: send a **SHA-256 id** instead of the document. **APQ** (`autoPersistQueries: true`): unknown id → server error → client retries with the full text (`useGETForPersistedQueryRetry` makes it a CDN-cacheable GET). **Persisted query list** (`operationManifest`): a *safelist* — the server runs only registered operations.
- Subscriptions**: `client.subscribe(subscription:)` → async sequence; over HTTP multipart, or WebSocket (`ApolloWebSocket`, `graphql-transport-ws`, 2.1+) with `SplitNetworkTransport` sending queries over HTTP and subscriptions over the socket.

REST or GraphQL? — say the trade-off

situation	lean
many clients, each needing different slices	GraphQL: one schema, per-screen queries
public, cache-heavy reads (CDN)	REST (or GraphQL GET + persisted ids)
simple CRUD, one client	REST: less machinery
large file upload / download	REST: or pre-signed URL; send the reference via GraphQL
live updates	subscriptions, or a plain WebSocket

Related: `http-deep` (caching, HTTP/2) · `websockets-realtime` (subscriptions' transport) · `ios-client-system-design` (normalised store, cursor pagination) · `urlsession-networking` · `api-design` (REST)