




In one line: ZFS is a volume manager and a filesystem in one: disks form **vdevs**, vdevs form a **pool**, and every **dataset** draws space from the whole pool. Nothing is overwritten in place — every change is written to **new blocks** and the tree of **checksummed** block pointers is re-rooted atomically, so a crash leaves the old or the new state, a **snapshot** is free, and silent corruption is *detected* and, given redundancy, *repaired*.



Copy-on-write: change one block



Write path



Vdevs — what each survives

vdev	survives	notes
mirror	all but 1 disk	fastest resilver (sequential resilver, 2.0); removable
raidz1/2/3	1 / 2 / 3 disks	distributed parity; widen by expansion (2.3)
druid1/2/3	1 / 2 / 3	raidz + distributed hot spare — fast rebuild (2.1)
special	= its own redundancy	metadata (+ small blocks); lose it, lose the pool — mirror it
log (SLOG)	ZIL content	low-latency SSD for sync writes; mirroring allowed
cache	—	L2ARC; persistent since 2.0; can't be mirrored
spare	—	hot spare taken over on a failure

Removal: spare, cache, log, mirror, single disk, special, dedup — only if **no raidz/draid** top-level vdev exists and all ashifts match. A raidz vdev can never be removed.

Integrity

- The checksum of every block lives in its **parent's** block pointer — a Merkle tree up to the uberblock, so a disk returning wrong data (bit rot, misdirected write) is caught on read. Default `fletcher4; sha256, sha512, skein, edonr, blake3` (2.2).
- With redundancy the bad copy is **rewritten from a good one** (self-healing). `zpool scrub` reads everything to find it before you need it.
- `copies=2|3` keeps extra copies even on one disk (charged to the dataset).
- `zpool status`: per-device **READ / WRITE / CKSUM** counters; files with permanent errors are listed by name.

Caching

ARC in RAM (adaptive: recently + frequently used). Default cap: the larger of **RAM** — **1 GiB** and **5/8 of RAM**. **L2ARC** on SSD extends it but needs RAM for its headers. Dirty data is capped at **RAM/10** by default.

Timeline

Sun starts ZFS 2001 · CDDL source in OpenSolaris 2005 · illumos fork 2010 · **OpenZFS** 2013 · 0.8 (2019) encryption, raw send, TRIM, device removal · **2.0** (2020) one codebase for Linux + FreeBSD, zstd · 2.1 (2021) dRAID · 2.2 (2023) block cloning, BLAKE3 · **2.3** (Jan 2025) RAIDZ expansion, fast dedup, Direct IO · 2.4 (Dec 2025). CDDL is held incompatible with the GPL, so on Linux ZFS ships as an out-of-tree module.

Datasets & properties (defaults)

recordsize	128 KiB (512 B–16 MiB); a <i>maximum</i> per block; new files only
volblocksize	zvols: 16 KiB since 2.2; fixed at creation
compression	on = lz4, default on since 2.2; also zstd (2.0), gzip
checksum	on = fletcher4
encryption	off; on = aes-256-gcm, per dataset, since 0.8
sync	standard · always · disabled (loses acknowledged writes)
atime	on, with relatime — set off for busy reads
dedup	off — every block's hash in a table that must stay hot

Workload advice (OpenZFS docs): InnoDB `recordsize=16K`, PostgreSQL 32K–128K, large sequential files 1M. **ashift** (sector size, 2^n) is fixed per vdev at creation — `ashift=12` for 4K disks, 12 or 13 for flash.

Snapshots, clones, replication

```
zpool create tank raidz2 sda sdb sdc sdd sde sdf
zfs create -o recordsize=1M tank/media
zfs snapshot tank/db@daily-1
zfs rollback tank/db@daily-1 # -r if newer snaps
zfs clone tank/db@daily-1 tank/t # writable, shared
zfs send -i @daily-1 tank/db@daily-2 |
  ssh backup zfs receive -u pool/db # delta only
zfs send -w tank/secret@s | ... # raw: still encrypted
zpool scrub tank; zpool status -x
```

Incremental send walks block birth times, not files — cost $\propto$ change, not size.

Traps

- ! A snapshot is not a backup: same pool, same disks — send it elsewhere.
- ! Adding a single disk as a top-level vdev to a raidz pool makes it the pool's weak point.
- ! RAIDZ expansion keeps old blocks at the old data:parity ratio until rewritten.
- ! A SLOG does nothing for async writes; L2ARC is no substitute for RAM.

Remember

Pool of vdevs · never overwrite · checksum in the parent · snapshot = old root · redundancy is per vdev.